

GenieKreator Configuration Spec: iLead Simulation

Oct 1, 2026 · @Manu Nair

Summary

In GenieKreator, every part of the iLead simulation that a learner sees, hears or is scored on is configurable data, except the engine itself. An author can turn the default Secure Capital Bank sales team into a pharma field force, a retail store or a dark store operation by changing context, people, images, voices, process, actions, events, scoring, gamification and report. They never touch code.

What is configurable vs locked

Configurable (12 areas, about 200 settings)	Locked engine (same for every client)
Organisation, scenario and product context	Turn loop: week, day budget, action, consequence
Branding, images, video and audio assets	State model maths: how deltas apply, clamps, decay
NPCs: roles, profiles, images, voices, personas, relationships	Evaluation pipeline: rubric, then outcome band, then authored consequences
Team structure, process stages, targets, KPIs	Safety guardrails for NPCs and participant inputs
Leadership model and fit rules	Data privacy and audit logging
Actions, live interactions, rubrics, consequences	Report formulas (inputs and weights are configurable)
Events and NPC initiated moments	
Time, pacing and play modes	
Gamification	
Report sections, skills framework, narratives	
Language, accessibility, delivery	
Governance: roles, approvals, versions	

How each setting gets its value

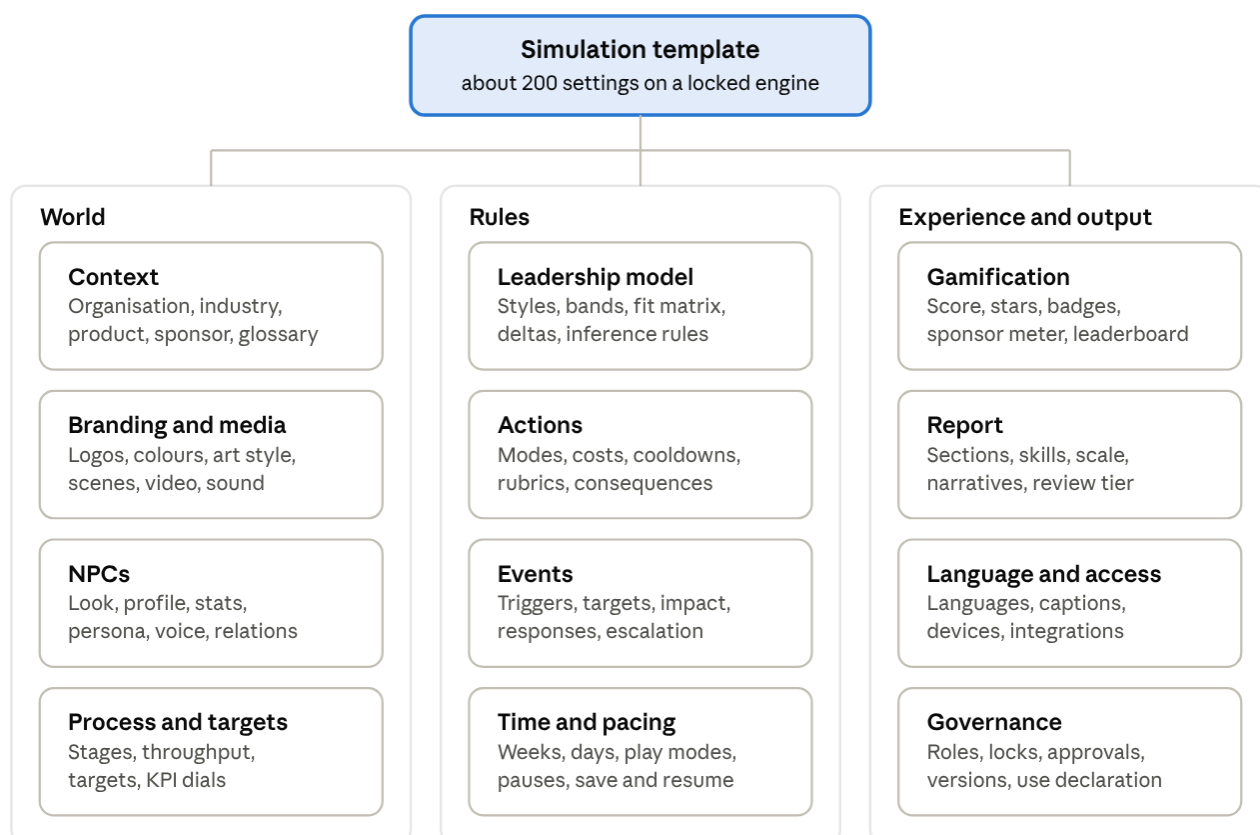
1. **AI generated from context.** The author gives a brief and optional uploads (org chart, process docs, values, product sheets, brand kit), and GenieKreator drafts a value for every setting.
2. **Author edited.** Every value can be changed in a form, or in plain language ("make the sponsor more demanding").
3. **Locked by admin.** A client admin can lock settings, such as brand colours or the skills framework, so authors cannot change them.

Each settings table below shows the setting, its type or options, the iLead default, and what the AI generates from context.

Configuration map

The 12 areas fall into three groups: the World the participant enters, the Rules that decide consequences, and the Experience and output around them.

Twelve configurable areas in three groups



configuration map · 12 areas in 3 groups

Areas are generated in this order, World first, so each later area can draw on the earlier ones.

Organisation, scenario and product context

This layer is the context engine input: everything else is generated from it. A richer brief produces more realistic NPCs, events and dialogue.

Organisation

Setting	Type or options	iLead default	AI generates from context
Organisation mode	Real client name, fictional twin, or fully fictional	Fictional (Secure Capital Bank)	Fictional twin that mirrors the client's industry and scale without naming it
Organisation name and tagline	Text	Secure Capital Bank	Name options in the client's naming style
Industry and sub industry	Picklist plus free text	Banking, retail lending	Taken from the brief
Size, structure, ownership	Picklists	Small, growth stage	Inferred from uploads
Region, city, office type	Text, picklist	Singapore City Branch	Local branch, store, plant or site name
Culture and values	Up to 6 values with one line each	Trusted client relationships	From the client's values document
Competitors	Up to 5 fictional or real names	Beta Bank, Hert Capital	Fictional rivals matching the market
Business situation	Turnaround, growth, change, crisis, merger, new launch	Turnaround after a weak manager	From the stated business challenge
Glossary	Term and meaning pairs	CASA, conversion	Client jargon pulled from uploads, used by NPCs
Policies NPCs follow	Upload or text (HR, compliance, sales conduct)	None	Rules for hiring, firing, rewards and claims

Product or service

Setting	Type or options	iLead default	AI generates from context
Product lines	1 to 5, name and description	CASA, Loans, Cards	From product sheets
Focus product	One of the above	Loans	The line tied to the business target
Value proposition and differentiators	Text	Low risk, attractive rate	Talking points NPCs and customers reference
Known weaknesses or criticisms	Text	Features not innovative	Seeds for events such as review site criticism
Customer segments	Up to 4	Individuals, companies	Segment personas for funnel and events
Unit of value	Revenue, units, accounts, NPS, cases closed, uptime	Revenue	From the client's main KPI
Currency and number format	Picklist	USD	From region

Scenario and role

Setting	Type or options	iLead default	AI generates from context
Participant role title	Text	Sales Director	Client's own job title (Area Manager, Store Manager, DHM)
Level and span	First line, mid, senior; team size	Mid manager, 10	From audience
Backstory	Text	Predecessor left the team in shatters	3 options matched to business situation
Sponsor persona	Name, title, image, voice, style (supportive, demanding, distant)	Roger Kent, CEO	Leader in the client's hierarchy (Regional Head, CHRO)

Setting	Type or options	iLead default	AI generates from context
Welcome message	Rich text	CEO welcome letter	Written in the sponsor's voice
Tone of the world	Formal, friendly, high pressure, playful	Professional	From culture
Realism level	Grounded, slightly dramatised, high drama	Grounded	Default grounded for senior audiences

Branding, images and media

Every visual can come from three sources: the client's upload, a licensed stock library, or AI generation in a chosen art style. One art style applies across all generated images, so the world looks consistent.

Brand and theme

Setting	Type or options	iLead default	AI generates from context
Brand ownership	KNOLSKAPE only, client only, both	KNOLSKAPE plus iLead	Client logo placement from brand kit
Logos	Upload: main, light, icon	iLead, KNOLSKAPE	
Colour theme	Primary, secondary, accent, status colours; light and dark	iLead blue and amber	Extracted from brand kit, checked for contrast
Fonts	Brand font upload or web font	iLead default	Closest web safe match
Simulation title and subtitle	Text	iLead DEMO	Client programme name
Loading and end screen art	Image	Office background	Generated in the art style

Images

Setting	Type or options	iLead default	AI generates from context
Art style	Photographic, illustrated flat, illustrated 3D, line art	Photographic	Recommended by audience and brand
Background scenes	Office, branch, store floor, warehouse, plant, hospital, field, virtual; per screen	Meeting room	Matched to the workplace in the brief
Workplace props	Uniforms, devices, signage, product packaging	None	Drawn from industry and product
NPC portraits	Upload, stock or generate; set per NPC (see NPCs)	Stock photos	Generated to match each profile
Expression set per NPC	Neutral, happy, concerned, frustrated, thinking	One photo	5 expressions generated from the base portrait, shown by mood
Event illustrations	Per event card	None	One image per event (CRM rollout, review site, festival rush)
Icons for actions and stages	Icon set picker	iLead icons	Matched to new stage names
Image safety	Auto checks for brand, offensive content, real person likeness		Blocks likeness of real people unless consented

Video and audio assets

Setting	Type or options	iLead default	AI generates from context
Sponsor welcome video	None, avatar video, uploaded video	Text letter	Avatar video of the sponsor reading the welcome
Intro and tutorial video	Upload or generated	Video link in HUD	Short walkthrough with the client's names
Ambient sound	Off, office, store, warehouse	Off	Matched to workplace

Setting	Type or options	iLead default	AI generates from context
UI sounds	Off, subtle, game like	Subtle	
Background music	Off, calm, upbeat	Off	

NPCs

An NPC is any AI character: team members, the sponsor, job candidates, customers, peers and HR. Each NPC is one record with 5 groups of settings (look, profile, stats, persona, voice), plus roster and relationship settings for the whole cast. The author can regenerate one group (for example only the voice) without touching the rest.

Roster level

Setting	Type or options	iLead default	AI generates from context
Team size	4 to 16	10	From span of control
Members per role	1 to 4, per stage	2	From the process
Supporting cast	Sponsor, HR partner, peer manager, customers, candidates	Sponsor only	Cast list matched to the actions switched on
Archetype mix	Required archetypes: top performer, low performer, complainer, role seeker, rival hire, quiet expert, new joiner, near retirement	5 archetypes implied	Balanced mix so every leadership style is needed at least twice
Diversity settings	Gender, age, ethnicity, nationality and name mix; or mirror the client's workforce data	Mixed	Mix matched to region with culturally correct names
Hiring pool	0 to 6 candidates with hidden true profiles	Not shown	Candidates with interview personas

Per NPC: identity and look

Setting	Type or options	iLead default	AI generates
Name, pronouns, age band	Text, picklist	Kent Goldberg	Regionally appropriate names
Portrait source	Upload, stock, generated	Stock photo	Generated in the art style
Look	Age, attire, setting, accessories	Office wear	Attire by industry (uniform, lab coat, store polo)
Expression set	5 moods	One photo	Generated from the portrait
Avatar mode	Static portrait, animated portrait with lip sync, video avatar	Static	Animated for 1:1 RolePlays

Per NPC: role and profile

Setting	Type or options	iLead default	AI generates
Role in team	Any stage or support role	Sales lead	From the process
Job title	Text	Implied	Client title
Previous company, tenure, experience	Text, numbers	The Japanese International Bank, 1 month, 3 years	Plausible history
Skills list	1 to 4 skills	Lead Generation, Negotiation	From role
Remarks visible to the leader	Text	Expected more, complaining	The visible cue for the archetype
Hidden concern	Text, revealed only through conversation	Implied	One concern per NPC tied to an event or action
Career goal	Text	Justin wants Conversion	Drives role fit and motivation

Per NPC: stats and fit

Setting	Type or options	iLead default	AI generates
Starting Skill, Morale, Result	0 to 100 each	Kent 30 / 22 / 49	Set by archetype, then balanced
Starting Trust	0 to 100	Not used	50, lower for wary archetypes
Role fit matrix	Skill, Motivation, Performance per role	Hidden	Derived from skills list and career goal
Needed style override	Optional fixed style per week	Derived from bands	Usually left to the band rules
Sensitivity weights	How strongly they react to recognition, criticism, change, workload	Uniform	Varied by persona (Jack very sensitive to recognition)

Per NPC: persona and behaviour

Setting	Type or options	iLead default	AI generates
Personality sliders	Openness, assertiveness, warmth, resilience, candour (1 to 5 each)	Implied in remarks	Set from archetype
Communication style	Direct, polite, verbose, terse, emotional, formal	Templated	Matched to culture
Attitude to the leader at start	Supportive, neutral, sceptical, hostile	Neutral	From backstory
Pressure response	Withdraws, argues, over promises, escalates		One per NPC
Opening lines and catchphrases	Text	"Hi, I'm Kent. How do you do?"	Written in persona
Knowledge scope	What the NPC knows: own work, team gossip, customer details, policy		Bounded to stop hallucination

Setting	Type or options	iLead default	AI generates
Off limits topics	List		Personal health details, politics, anything in client policy
Memory depth	Last 3, 5 or all interactions		5 by default

Per NPC: voice

Setting	Type or options	iLead default	AI generates
Voice source	Library voice, client licensed voice, none (text only)	None	Library voice matching age, gender and region
Language and accent	Any supported language and regional accent		Accent of the client region, mixed if the team is global
Pitch, pace, warmth	Sliders		From persona
Emotional range	Flat, moderate, expressive; mood linked		Moves with the NPC's current Morale
Pronunciation dictionary	Word and phonetic pairs		Names, product terms, acronyms
Voice consent record	Required for any cloned or licensed voice		Blocks unlicensed cloning

Relationships between NPCs

Setting	Type or options	iLead default	AI generates
Relationship map	Pairs with type: allies, rivals, mentor, friends, conflict	None	3 to 5 links that create ripples
Influence	Who sways team opinion (0 to 100)	Implied (Kent complains to others)	Higher for top performers and complainers

Setting	Type or options	iLead default	AI generates
Ripple rules	When A is rewarded, criticised or moved, how B reacts	Jack reacts to others' bonuses	Generated from the map and sensitivity weights

Team structure, process, targets and KPIs

The sales funnel is one instance of a general process pipeline. Any chain of stages where work flows from one role to the next can replace it: a store, a dark store, a contact centre, a field force, a project team.

Process pipeline

Setting	Type or options	iLead default	AI generates from context
Process type	Sales funnel, service flow, operations line, project delivery, custom	Sales funnel	From the client's process docs
Number of stages	3 to 7	5	From the process
Stage names, icons, tooltips	Text, icon, rich text	Sales lead, Qualify, Proposal, Negotiate, Conversion	Client's own stage names (Inbound, Pick, Pack, Dispatch, Deliver)
Work unit	Lead, order, ticket, patient, case, task	Lead	From the process
Ideal throughput per stage per week	Numbers	3, 2, 1, 1, 1	Back calculated from the target
Throughput formula weights	How Skill, Morale and Result of the 2 stage owners drive output	Engine default	Balanced in the quality gate
Bottleneck rules	Stage capacity, carry over of unprocessed work	Carry over on	

Setting	Type or options	iLead default	AI generates from context
Parallel or branching stages	Optional: a stage can split (for example returns vs deliveries)	Linear	Suggested only when the process branches

Targets and KPIs

Setting	Type or options	iLead default	AI generates from context
Primary target	Revenue, units, SLA %, NPS, cases resolved	\$240,000	From the client's main KPI
Value per completed unit	Number	\$30,000 per conversion	Realistic for the industry
Secondary targets	Up to 3: team skill, morale, attrition, quality	Improve skill, morale, performance	From the brief
Team KPI dials	Choose 3 to 5 from Skill, Morale, Result, Trust, Engagement, Quality, Safety	Skill, Morale, Result	Adds the client's quality KPI if relevant
KPI labels	Rename any dial		Client language (Productivity instead of Result)
Weekly drift	Automatic decline of morale or result when nothing is done	About 3 points per week	
Win condition	Hit target, hit target and keep morale above X, or best score	Target	

Leadership model and scoring rules

The leadership model is the yardstick for every weekly choice and every live interaction. Authors can keep the iLead model, pick another framework, or load the client's own.

Setting	Type or options	iLead default	AI generates from context
Framework	iLead 4 styles, Situational Leadership style labels, coaching styles, client framework upload	Directing, Guiding, Partnering, Entrusting	Maps a client framework onto the engine's 4 slots
Number of styles	3 to 6	4	
Style names and definitions	Text per style	In game definitions	Client terms (for example Tell, Sell, Share, Delegate)
Readiness bands	Thresholds on current Skill and Morale (and optionally Trust)	Hidden engine values	Proposed bands, checked in the balance test
Fit matrix	Which style suits which band	Hidden	Drafted from the framework
Per member exceptions	Style override for a person and week	None	Used for special stories (an NPC in crisis needs Guiding)
Role misfit penalty	How much a poor role fit reduces style gains	Strong (Justin failed all styles)	Moderate by default, so feedback stays readable
Delta table	Morale and Result change for match and mismatch	About +2 to +3, -1 to -2	
Team feedback messages	Copy by share of matches (below half, half, majority, all)	3 messages seen	Rewritten in the sponsor's voice
Style tags for action options	Which style each static option and each live behaviour counts as	Fixed per option	Generated with the option text

Setting	Type or options	iLead default	AI generates from context
Style inference rules for live interactions	Behaviour indicators per style (for example detailed instructions = Directing, asking for their plan = Entrusting)	Not in current build	Indicators drafted per framework
Intent vs action rule	How often a gap costs Trust	Not in current build	Gap 2 weeks in a row costs Trust

Actions and live interactions

The action catalogue starts with the 13 iLead actions and the 2 new live ones (Reply to NPC messages, Sponsor briefing). Authors can rename, hide, clone or create actions. Each action has a common record; live actions add an interaction record.

Per action

Setting	Type or options	iLead default	AI generates from context
Name, icon, description	Text, icon	Send email, Coach member	Client language ("Huddle" instead of Meet the team)
Scope	Team, individual, multi select	Per action	
Mode	Static, hybrid, live	See iLead 2.0 design	Recommended by action type
Day cost	0 to 3	1 or 2	
Cooldown	Days before reuse	Energize: 10 days	
Limits	Max targets, max uses per week or run	Email 3, training 3	
Eligibility rules	Role coverage, minimum one per role, member available	As in iLead	

Setting	Type or options	iLead default	AI generates from context
Prerequisites	Action that should come first, with a penalty if skipped	Assess before Swap	Suggested chains (Assess then Swap then Train)
Options (static)	2 to 4 options, each with text and a style tag	4 per action	Rewritten in client context
Outcome quotes (static)	Reply per option and band	Templated	Varied quotes in each NPC's persona
Base consequences	Skill, Morale, Result, Trust deltas for target and ripple	Engine values	Balanced in the quality gate
Unlock condition	Available from week N, or after an event or sponsor unlock	All available	

Per live interaction

Setting	Type or options	iLead default	AI generates from context
Format	Email, chat, 1:1 RolePlay, team meeting, sponsor briefing, interview, written plan	Per action	
Input modes allowed	Text, audio, both	Both	
Time limit	Minutes, soft or hard stop	3 to 5	
Turn limit	Max participant turns	12 for RolePlays	
Participant brief	What the participant sees before starting: goal, context, NPC card		Written from the scenario and NPC state
NPC brief	What the NPC wants, fears, hides and will accept		Drawn from persona and hidden concern
Difficulty	Easy, standard, tough NPC	Standard	Tougher for senior audiences

Setting	Type or options	iLead default	AI generates from context
Opening	Who speaks first, opening line	NPC	
Rubric dimensions	2 to 4 skills with behavioural anchors per band		Drafted from the linkage matrix
Concern detection	Phrases or ideas that count as surfacing the hidden concern		Generated with paraphrase variants
Red flags	Behaviours that force the Harmful band (blame, discrimination, policy breach)		From client policy
Outcome bands	Strong, Adequate, Weak, Harmful; band names editable	4 bands	
Consequence table	Deltas per band for target, bystanders and sponsor		Balanced in the quality gate
Triggers per band	Follow up events, promises logged, concern resolved or escalated		Linked to the event deck
Hints	Off, on request, after a Weak band	On request	Coaching tips in the client's language
Calibration set	6 to 12 sample answers with author labels		Generated samples for the author to label
Email specifics	Recipients allowed, CC and BCC, attachments, reply all reactions		
Meeting specifics	Attendees, agenda required, NPC to NPC talk on or off		
Interview specifics	Number of candidates, question bank on or off, hidden true profile	2 candidates	

Events and NPC initiated moments

Events make the world move without the participant. GenieKreator generates an 8 week deck from the context (a CRM rollout becomes a WMS rollout for a warehouse) and the author reorders, edits or adds cards.

Setting	Type or options	iLead default	AI generates from context
Event type	Impact, signal, capacity, diagnostic, opportunity, crisis	4 types seen	Mix balanced across weeks
Theme	Change, competitor, reputation, personal, attrition, compliance, customer, market	CRM, review site, personal, leave	Themes from the client's real pressures
Trigger	Fixed week and day, random within a window, conditional (state or action based)	Week start or random	Conditional triggers such as "Morale below 30 for 2 weeks"
Probability	0 to 100% for random events		
Target	Team, role, named NPC, sponsor	Team, role, member	
Delivery	News bulletin card, modal, NPC chat message, email in inbox, sponsor call	Bulletin, modal	Delivery matched to the event
Text and image	Rich text, image	Short card	Written in the world's tone; illustration generated
Immediate impact	Deltas, capacity loss (days away), funnel change	Seen per event	
Expected response	Action or behaviour the engine rewards	Implied	For example: reply to Kent within 2 days with support
Response window	Days before an ignored event escalates		2 to 5 days

Setting	Type or options	iLead default	AI generates from context
Escalation chain	Follow up event if ignored or mishandled		Complaint to sponsor, transfer request
Repeat rules	Once, recurring, cooldown	Once	
Label on card	Show or hide "no impact on result" style labels	Shown	Hidden by default, since labels can mislead
Deck order and pacing	Events per week, intensity curve	1 to 3 per week	Rising pressure to Week 6, then a recovery window

Time, pacing and play modes

Setting	Type or options	iLead default	AI generates from context
Play mode	Full, Standard, Lite, custom	Full in iLead 2.0	From session length in the brief
Simulated weeks	2 to 12	8	
Days per week	3 to 7	5	Matches the client's shift pattern (6 days for retail)
Time unit label	Week, sprint, shift, month	Week	
Live interaction cap	Per week and per run	2 per week	
Real time limit	Off, total minutes, minutes per week	20 minutes	Off or generous in Full mode
Clock pause rules	Pause during live interactions, modals, reports	No pauses	Pause during live interactions
Save and resume	On or off; resume window in days	Off	On for Full mode
Sittings	1 to 4, with week breaks defined	1	2 sittings for Full

Setting	Type or options	iLead default	AI generates from context
Self paced vs facilitated	Self paced, facilitator controls week advance, cohort sync	Self paced	
Facilitator controls	Pause cohort, inject event, extend time, reset a participant	None	
Practice week	Optional Week 0 tutorial with no scoring	Guided tour only	

Gamification settings

Score logic is defined in the [iLead 2.0 design doc](#). Every element below can be switched off, retuned or rethemed.

Setting	Type or options	iLead 2.0 default	AI generates from context
Elements on or off	Score, stars, streaks, badges, sponsor meter, unlocks, Team Pulse, leaderboard, tiers	All on	Off set for selection use
Score weights	Business, People, Leadership (sum 100)	30 / 30 / 40	From the client's stated outcome
Score scale	0 to 100, 0 to 1000, custom	0 to 1000	
Star thresholds	3 values	50 / 70 / 85	
Streak rules	Length to start, bonus per week, cap	3 weeks, +25, 100	
Badge library	Add, edit, remove; rule builder (event, condition, count); name, icon, copy	10 badges	Names and icons in the client's theme
Sponsor meter	Start value, gains and losses per trigger	50; per design doc	
Unlock rewards	Bonus day, hire budget, extra team activity, custom	3 unlocks	Rewards that fit the scenario

Setting	Type or options	iLead 2.0 default	AI generates from context
Tier names and thresholds	2 to 5 tiers	Bronze to Platinum	Client themed names
Leaderboard	Scope (cohort, business unit, global), size, anonymity, window	Cohort, top 10, named	Anonymous for mixed seniority cohorts
Celebrations	Animation level: none, subtle, full	Subtle	

Report settings

Setting	Type or options	iLead 2.0 default	AI generates from context
Report sections	Choose and order the 10 sections	All 10	Shorter set for Lite mode
Skills framework	iLead default 8 skills, KNOLSKAPE Skills Ontology, client framework upload	iLead default	Maps client skills to evidence sources
Linkage matrix	Skills by interaction, minimum 2 observations per skill	Default matrix	Rebuilt for the chosen actions; warns on any skill under 2
Rating scale	Levels, labels, colours	Novice to Role Model, 5 levels	
Behavioural anchors	Text per skill and level	Drafted	From the client's skill definitions
Narrative bank	Copy per band, style and section	Template copy	Written in the client's tone and language
Evidence quotes	Number per skill, redaction of names	2 per skill	
Development plan	Practice activity, on the job action, check in date per priority	3 priorities	Linked to the client's learning catalogue

Setting	Type or options	iLead 2.0 default	AI generates from context
Audiences	Participant, manager, L&D admin, cohort view	Participant and L&D	
Delivery	In app, PDF, email, LMS record	In app and PDF	
Branding	Same as simulation, or report specific	Same	
Human review tier	Off, sample audit, full assessor review	Sample audit	Full review suggested for high stakes use

Language, accessibility, delivery and governance

Setting	Type or options	Default	AI generates from context
Content languages	One or many; participant picks at start	English	Full translation of copy, NPC dialogue and report
Mixed language play	Participant may speak or type in a different language from the UI	Off	
Locale formats	Dates, numbers, currency, names order	From region	
Accessibility	Captions for all voice, transcripts, keyboard only play, screen reader labels, reduced motion, high contrast, text size	All on	
Typed fallback in voice interactions	Always available	On	
Devices	Desktop, tablet, mobile (Lite only)	Desktop and tablet	
Low bandwidth mode	Static portraits, text NPCs, no ambient audio	Auto	

Setting	Type or options	Default	AI generates from context
Integration	SCORM 1.2, SCORM 2004, xAPI, LTI, API; SSO	Platform default	
Data and consent	Audio consent screen, retention period, transcript storage, region of storage	Consent on, retention set by client	
Roles and permissions	Admin, author, reviewer, facilitator, viewer	All roles	
Locks	Admin can lock any setting for authors	Brand, skills framework	
Approval workflow	Draft, in review, approved, published; reviewer sign off per area	On	
Versioning	Version history, compare, roll back, clone to a new client	On	
Templates	Save any configured simulation as a reusable template	On	
Use declaration	Development or selection use (changes report, leaderboard and review defaults)	Development	

Authoring workflow, AI assist and quality gates

The brief (about 8 questions)

1. Who is the participant: role title, level, team size, region, language?
2. What industry and organisation, real or fictional?
3. What product or service, and what is the main business target?
4. What is the business situation: turnaround, growth, change, crisis?

5. Which 3 to 6 skills matter most?
6. How long is the session, and is it facilitated or self paced?
7. Is it for development or for selection decisions?
8. Uploads: org chart, process docs, values, product sheets, brand kit, skills framework (all optional).

Generation order and editing

1. Context and scenario, then process and targets, then NPCs, then leadership model, then actions and interactions, then events, then gamification, then report. Each later step reads the earlier ones, so a changed process regenerates stage specific events and rubrics.
2. Regenerate at any scope: the whole simulation, one area, one NPC, or one field ("new voice for Peter").
3. Plain language edits: "make the sponsor tougher", "add a night shift lead who resists change", "set it in Dubai".
4. Change impact warnings: editing stats, deltas or thresholds marks the balance test as stale until it is rerun.

Quality gates before publish

Gate	What it checks	Pass rule
Balance test	AI players (strong, one style, careless, random) play 50 or more runs	Strong reaches the target and Gold tier; careless stays below target and in Bronze
Style coverage	Every style is the right answer for at least 2 members across the run	No unused or dominant style
Rubric calibration	AI bands vs author labels on the calibration set	At least 85% agreement per interaction
Persona test	Each NPC faces off topic, hostile and manipulative inputs	Stays in role, follows off limits rules
Content safety	Images, names, voices, copy	No real person likeness or unlicensed voice; no flagged content
Linkage check	Report skills have enough evidence	Each skill has at least 2 observation points
Accessibility check	Contrast, captions, keyboard flow	Meets WCAG 2.2 AA

Gate	What it checks	Pass rule
Preview	Author plays one full week with a debug panel showing hidden state	Author sign off